
MODULE *DistributedLock*

EXTENDS *Naturals, FiniteSets, Sequences, TLC*

The set of clients

CONSTANT *Clients*

Client states

CONSTANTS *Active, Inactive*

Message types

CONSTANT *LockRequest, LockResponse, TryLockRequest, TryLockResponse, UnlockRequest, UnlockResponse*

An empty constant

CONSTANT *Nil*

The current lock holder

VARIABLE *lock*

The lock queue

VARIABLE *queue*

The current lock *ID*

VARIABLE *id*

Session states

VARIABLE *sessions*

$serverVars \triangleq \langle lock, id, queue, sessions \rangle$

Client states

VARIABLE *clients*

$clientVars \triangleq \langle clients \rangle$

Client requests

VARIABLE *requests*

Server responses

VARIABLE *responses*

Variable to track the total number of messages sent for use in state constraints when model checking

VARIABLE *messageCount*

$messageVars \triangleq \langle requests, responses, messageCount \rangle$

The invariant checks that:

- * No client can hold more than one lock at a time
- * No two clients hold a lock with the same *ID*

* The lock is held by an active session

Note that more than one client may believe itself to hold the lock at the same time, *e.g.* if a client's session has expired but the client hasn't been notified, but lock *IDs* must be unique and monotonically increasing.

$TypeInvariant \triangleq$

$\wedge \forall c \in \text{DOMAIN } clients : Cardinality(clients[c].locks) \in 0 .. 1$

$\wedge \forall c1, c2 \in \text{DOMAIN } clients : c1 \neq c2 \Rightarrow Cardinality(clients[c1].locks \cap clients[c2].locks) = 0$

$\wedge lock \neq Nil \Rightarrow sessions[lock.client].state = Active$

Returns a sequence with the head removed

$Pop(q) \triangleq SubSeq(q, 2, Len(q))$

Sends a message on the given client's channel

$Send(m, c) \triangleq$

$\wedge requests' = [requests \text{ EXCEPT } ![c] = Append(requests[c], m)]$

$\wedge messageCount' = messageCount + 1$

$\wedge \text{UNCHANGED } \langle responses \rangle$

Removes a message from the given client's channel

$AcceptRequest(m, c) \triangleq$

$\wedge requests' = [requests \text{ EXCEPT } ![c] = Pop(requests[c])]$

$\wedge messageCount' = messageCount + 1$

$\wedge \text{UNCHANGED } \langle responses \rangle$

Removes a message from the given server's channel

$AcceptResponse(m, c) \triangleq$

$\wedge responses' = [responses \text{ EXCEPT } ![c] = Pop(responses[c])]$

$\wedge messageCount' = messageCount + 1$

$\wedge \text{UNCHANGED } \langle requests \rangle$

Removes the last message from the client's channel and appends a message to the given server's channel

$Reply(m, c) \triangleq$

$\wedge requests' = [requests \text{ EXCEPT } ![c] = Pop(requests[c])]$

$\wedge responses' = [responses \text{ EXCEPT } ![c] = Append(responses[c], m)]$

$\wedge messageCount' = messageCount + 1$

Handles a lock request. If the lock is not currently held by another process, the lock is granted to the client. If the lock is held by a process, the request is added to a queue.

$HandleLockRequest(m, c) \triangleq$

$\vee \wedge sessions[c].state \neq Active$

$\wedge AcceptRequest(m, c)$

$\wedge \text{UNCHANGED } \langle clientVars, serverVars \rangle$

$\vee \wedge sessions[c].state = Active$

$\wedge lock = Nil$

$$\begin{aligned}
& \wedge lock' = m @@"client" :> c) \\
& \wedge id' = id + 1 \\
& \wedge Reply([type \mapsto LockResponse, acquired \mapsto TRUE, id \mapsto id'], c) \\
& \wedge UNCHANGED \langle queue, sessions, clientVars \rangle \\
\vee & \wedge sessions[c].state = Active \\
& \wedge lock \neq Nil \\
& \wedge queue' = Append(queue, m @@"client" :> c)) \\
& \wedge AcceptRequest(m, c) \\
& \wedge UNCHANGED \langle lock, id, sessions, clientVars \rangle
\end{aligned}$$

Handles a *tryLock* request. If the lock is not currently held by another process, the lock is granted to the client. Otherwise, the request is rejected.

$$\begin{aligned}
HandleTryLockRequest(m, c) & \triangleq \\
& \vee \wedge sessions[c].state \neq Active \\
& \quad \wedge AcceptRequest(m, c) \\
& \quad \wedge UNCHANGED \langle clientVars, serverVars \rangle \\
& \vee \wedge sessions[c].state = Active \\
& \quad \wedge lock = Nil \\
& \quad \wedge lock' = m @@"client" :> c) \\
& \quad \wedge id' = id + 1 \\
& \quad \wedge Reply([type \mapsto LockResponse, acquired \mapsto TRUE, id \mapsto id'], c) \\
& \quad \wedge UNCHANGED \langle queue, sessions, clientVars \rangle \\
& \vee \wedge sessions[c].state = Active \\
& \quad \wedge lock \neq Nil \\
& \quad \wedge Reply([type \mapsto LockResponse, acquired \mapsto FALSE], c) \\
& \quad \wedge UNCHANGED \langle clientVars, serverVars \rangle
\end{aligned}$$

Handles an unlock request. If the lock is currently held by the given client, it will be unlocked. If any client's requests are pending in the queue, the next lock request will be removed from the queue and the lock will be granted to the requesting client.

$$\begin{aligned}
HandleUnlockRequest(m, c) & \triangleq \\
& \vee \wedge sessions[c].state \neq Active \\
& \quad \wedge AcceptRequest(m, c) \\
& \quad \wedge UNCHANGED \langle clientVars, serverVars \rangle \\
& \vee \wedge sessions[c].state = Active \\
& \quad \wedge lock = Nil \\
& \quad \wedge AcceptRequest(m, c) \\
& \quad \wedge UNCHANGED \langle clientVars, serverVars \rangle \\
& \vee \wedge sessions[c].state = Active \\
& \quad \wedge lock \neq Nil \\
& \quad \wedge lock.client = c \\
& \quad \wedge lock.id = m.id \\
& \quad \wedge \vee \wedge Len(queue) > 0 \\
& \quad \quad \wedge LET next \triangleq Head(queue) \\
& \quad \quad IN \\
& \quad \quad \wedge lock' = next
\end{aligned}$$

Sends a lock request to the cluster with a unique *ID* for the client.

$$\begin{aligned} \text{Lock}(c) \triangleq & \\ & \wedge \text{clients}[c].\text{state} = \text{Active} \\ & \wedge \text{Send}([type \mapsto \text{LockRequest}, id \mapsto \text{clients}[c].\text{next}], c) \\ & \wedge \text{clients}' = [\text{clients} \text{ EXCEPT } ![c].\text{next} = \text{clients}[c].\text{next} + 1] \\ & \wedge \text{UNCHANGED } \langle \text{serverVars} \rangle \end{aligned}$$

Sends a try lock request to the cluster with a unique *ID* for the client.

$$\begin{aligned} \text{TryLock}(c) \triangleq & \\ & \wedge \text{clients}[c].\text{state} = \text{Active} \\ & \wedge \text{Send}([type \mapsto \text{TryLockRequest}, id \mapsto \text{clients}[c].\text{next}], c) \\ & \wedge \text{clients}' = [\text{clients} \text{ EXCEPT } ![c].\text{next} = \text{clients}[c].\text{next} + 1] \\ & \wedge \text{UNCHANGED } \langle \text{serverVars} \rangle \end{aligned}$$

Sends an unlock request to the cluster if the client is active and current holds a lock.

$$\begin{aligned} \text{Unlock}(c) \triangleq & \\ & \wedge \text{clients}[c].\text{state} = \text{Active} \\ & \wedge \text{Cardinality}(\text{clients}[c].\text{locks}) > 0 \\ & \wedge \text{Send}([type \mapsto \text{UnlockRequest}, id \mapsto \text{CHOOSE } l \in \text{clients}[c].\text{locks} : \text{TRUE}], c) \\ & \wedge \text{clients}' = [\text{clients} \text{ EXCEPT } ![c].\text{locks} = \text{clients}[c].\text{locks} \setminus \{\text{CHOOSE } l \in \text{clients}[c].\text{locks} : \text{TRUE}\}] \\ & \wedge \text{UNCHANGED } \langle \text{serverVars} \rangle \end{aligned}$$

Handles a lock response from the cluster. If the client's session is expired, the response is ignored. If the lock was acquired successfully, it's added to the client's lock set.

$$\begin{aligned} \text{HandleLockResponse}(m, c) \triangleq & \\ & \wedge \vee \wedge \text{clients}[c].\text{state} = \text{Inactive} \\ & \quad \wedge \text{UNCHANGED } \langle \text{clientVars}, \text{serverVars} \rangle \\ & \vee \wedge \text{clients}[c].\text{state} = \text{Active} \\ & \quad \wedge m.\text{acquired} \\ & \quad \wedge \text{clients}' = [\text{clients} \text{ EXCEPT } ![c].\text{locks} = \text{clients}[c].\text{locks} \cup \{m.\text{id}\}] \\ & \quad \wedge \text{UNCHANGED } \langle \text{serverVars} \rangle \\ & \vee \wedge \text{clients}[c].\text{state} = \text{Active} \\ & \quad \wedge \neg m.\text{acquired} \\ & \quad \wedge \text{UNCHANGED } \langle \text{clientVars}, \text{serverVars} \rangle \\ & \wedge \text{AcceptResponse}(m, c) \end{aligned}$$

Receives a message from/to the given client from the head of the client's message queue.

$$\begin{aligned} \text{Receive}(c) \triangleq & \\ & \vee \wedge \text{Len}(\text{requests}[c]) > 0 \\ & \quad \wedge \text{LET } \text{message} \triangleq \text{Head}(\text{requests}[c]) \\ & \quad \text{IN} \\ & \quad \vee \wedge \text{message.type} = \text{LockRequest} \end{aligned}$$

$$\begin{aligned}
& \wedge \text{HandleLockRequest}(\text{message}, c) \\
\vee & \wedge \text{message.type} = \text{TryLockRequest} \\
& \wedge \text{HandleTryLockRequest}(\text{message}, c) \\
\vee & \wedge \text{message.type} = \text{UnlockRequest} \\
& \wedge \text{HandleUnlockRequest}(\text{message}, c) \\
\vee & \wedge \text{Len}(\text{responses}[c]) > 0 \\
& \wedge \text{LET } \text{message} \triangleq \text{Head}(\text{responses}[c]) \\
& \text{IN} \\
& \vee \wedge \text{message.type} = \text{LockResponse} \\
& \wedge \text{HandleLockResponse}(\text{message}, c)
\end{aligned}$$

Initial state predicate

$$\begin{aligned}
\text{Init} & \triangleq \\
& \wedge \text{requests} = [c \in \text{Clients} \mapsto \langle \rangle] \\
& \wedge \text{responses} = [c \in \text{Clients} \mapsto \langle \rangle] \\
& \wedge \text{messageCount} = 0 \\
& \wedge \text{lock} = \text{Nil} \\
& \wedge \text{queue} = \langle \rangle \\
& \wedge \text{id} = 0 \\
& \wedge \text{clients} = [c \in \text{Clients} \mapsto [\text{state} \mapsto \text{Active}, \text{locks} \mapsto \{\}, \text{next} \mapsto 1]] \\
& \wedge \text{sessions} = [c \in \text{Clients} \mapsto [\text{state} \mapsto \text{Active}]]
\end{aligned}$$

Next state predicate

$$\begin{aligned}
\text{Next} & \triangleq \\
& \vee \exists c \in \text{DOMAIN } \text{clients} : \text{Receive}(c) \\
& \vee \exists c \in \text{DOMAIN } \text{clients} : \text{Lock}(c) \\
& \vee \exists c \in \text{DOMAIN } \text{clients} : \text{TryLock}(c) \\
& \vee \exists c \in \text{DOMAIN } \text{clients} : \text{Unlock}(c) \\
& \vee \exists c \in \text{DOMAIN } \text{clients} : \text{ExpireSession}(c) \\
& \vee \exists c \in \text{DOMAIN } \text{clients} : \text{CloseSession}(c)
\end{aligned}$$

The specification includes the initial state predicate and the next state

$$\text{Spec} \triangleq \text{Init} \wedge \square[\text{Next}]_{\langle \text{serverVars}, \text{clientVars}, \text{messageVars} \rangle}$$

\ * Modification History
\ * Last modified Sun Jan 28 10:20:37 PST 2018 by jordanhalterman
\ * Created Fri Jan 26 13:12:01 PST 2018 by jordanhalterman